

Using pytorch models in c++ projects

Georgy Derevyanko

Plan

Pytorch C++ API intro:

- **TH library: THTensor, THStorage**
- **ATen: Tensor, Storage, Array ...**

Converting model:

- **JIT**
- **Parsing JIT in C++**
- **Compiling using CMake and SConstruct**

Showing that it works:

- **Demo**

TH Library: moved to c10 namespace

Old code:

<https://github.com/pytorch/pytorch/tree/master/aten/src/TH/generic>

Code like this (C interface):

```
THTensor_(nDimension)(const THTensor *self);
```

is converted to:

```
THFloatTensor_nDimension(const THTensor *self);
```

```
THDoubleTensor_nDimension(const THTensor *self);
```

```
THByteTensor_nDimension(const THTensor *self);
```

...

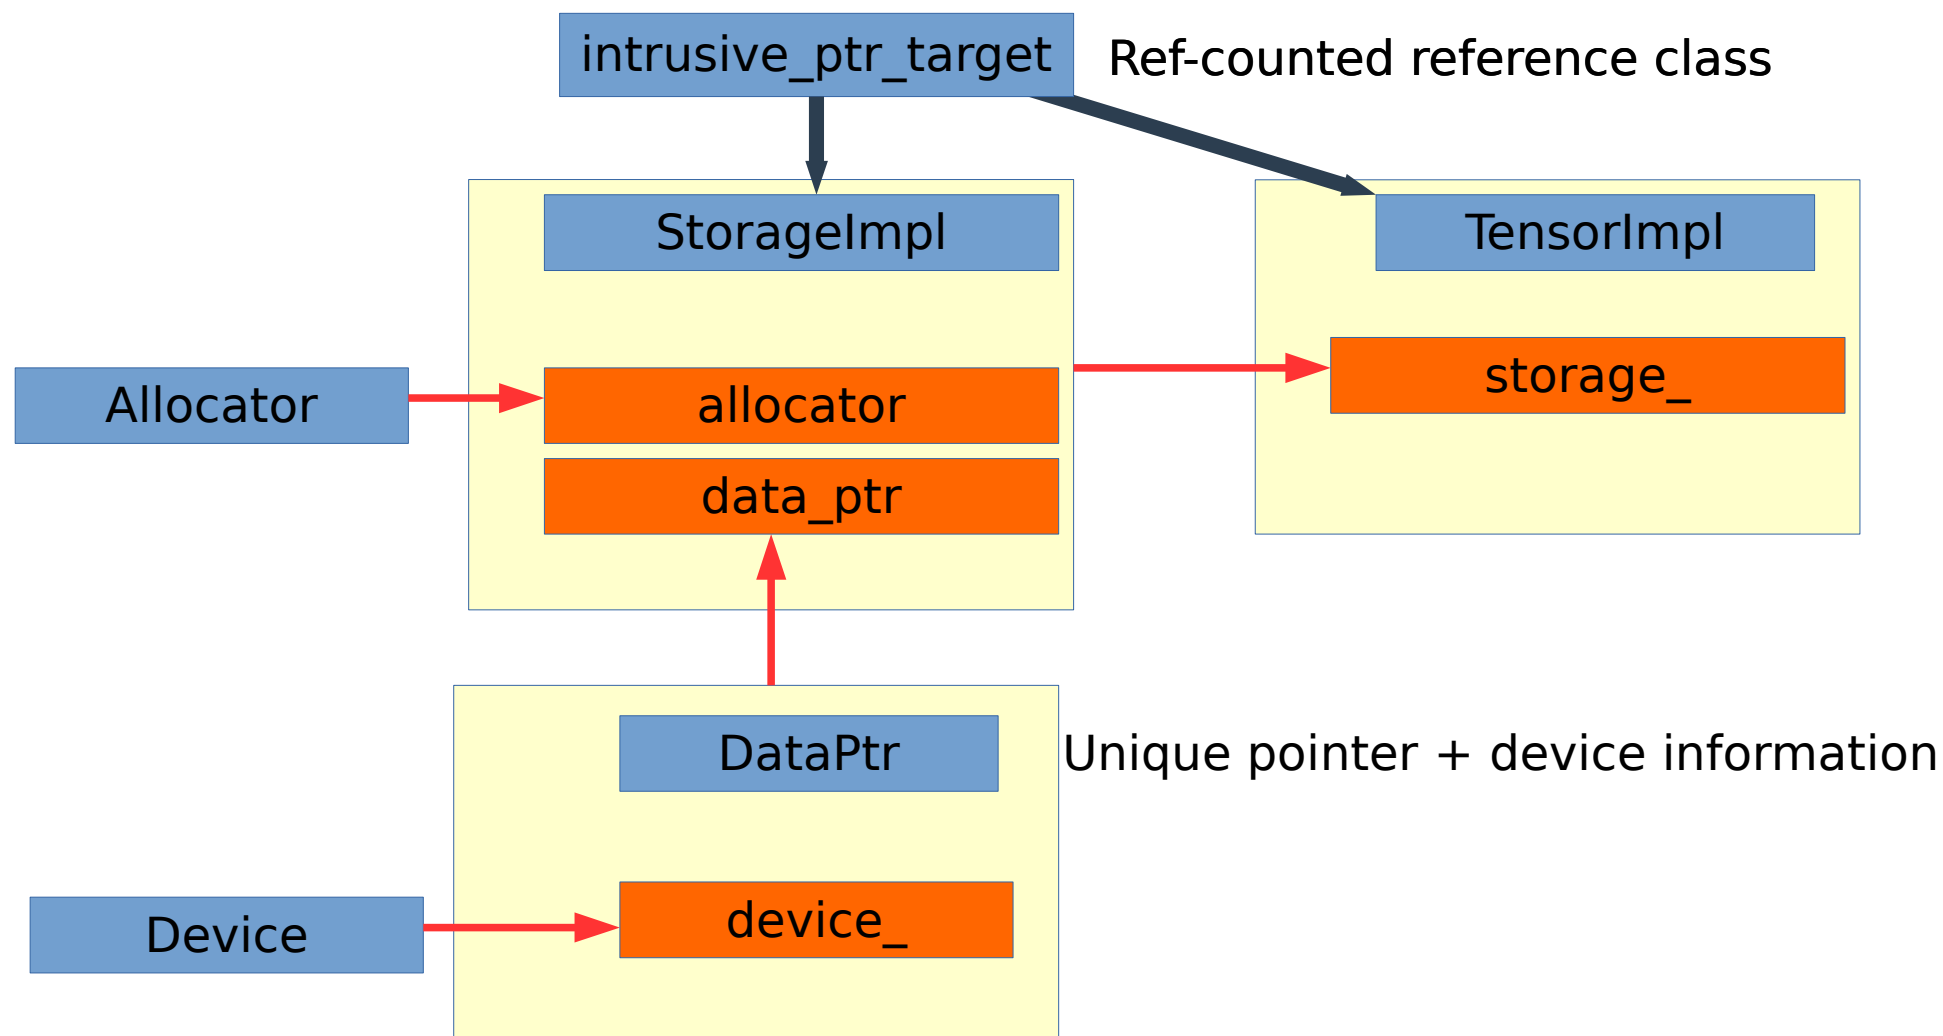
Tensor properties, creation, manipulation: `THTensor.h`, `THTensor.cpp`

Storage: `THStorage.h`, `THStorage.cpp`

...

THLibrary is only an interface to `c10:TensorImpl`

c10: general architecture



ATen: general architecture

c10::Storage wraps c10::StorageImpl

...

C10Tensor wraps c10::TensorImpl

**at::Tensor wraps pointer to the tensor
implementation**

c10 + at: Arrays and TensorList

```
std::vector<torch::Tensor> cc2({x, torch::zeros({batch_size, x.size(1), 1, 1}, this_device)});
```

```
torch::Tensor y2 = at::cat(at::TensorList(cc2), 2);
```



<https://github.com/pytorch/pytorch/blob/master/c10/util/ArrayRef.h>

in <https://github.com/pytorch/pytorch/blob/master/aten/src/ATen/core/Tensor.h>:

```
using TensorList = ArrayRef<Tensor>;
```

Let's say you've trained your model in python.

```
model.eval()
```

```
dummy_input = torch.ones(1,1024, dtype=torch.float, device='cpu')
```

```
output = model(dummy_input)
```

```
trace = torch.jit.trace(model, dummy_input)
```

```
trace.save(output_filename)
```

Documentation (running traced script**):**

<https://pytorch.org/docs/stable/jit.html>

https://pytorch.org/tutorials/advanced/cpp_export.html

Reimplementing the model in C++

```
struct CREPE : torch::nn::Module{
    torch::Tensor pad1;
    int batch_size;
    at::TensorOptions this_device;
    CREPE(at::TensorOptions device, int capacity=32)
    : conv1(torch::nn::Conv2dOptions(1, capacity*32, /*kernel_size=*/{512, 1}).stride({4,1}).padding({254,0})),
      conv1_bn(torch::nn::BatchNormOptions(capacity*32)),
      conv2(torch::nn::Conv2dOptions(capacity*32, capacity*4, {64, 1}).padding({31,0})),
      conv2_bn(torch::nn::BatchNormOptions(capacity*4)),
      conv3(torch::nn::Conv2dOptions(capacity*4, capacity*4, {64, 1}).padding({31,0})),
      conv3_bn(torch::nn::BatchNormOptions(capacity*4)),
      conv4(torch::nn::Conv2dOptions(capacity*4, capacity*4, {64, 1}).padding({31,0})),
      conv4_bn(torch::nn::BatchNormOptions(capacity*4)),
      conv5(torch::nn::Conv2dOptions(capacity*4, capacity*8, {64, 1}).padding({31,0})),
      conv5_bn(torch::nn::BatchNormOptions(capacity*8)),
      conv6(torch::nn::Conv2dOptions(capacity*8, capacity*16, {64, 1}).padding({31,0})),
      conv6_bn(torch::nn::BatchNormOptions(capacity*16)),
      fc(torch::nn::LinearOptions(capacity*64, 360)),
      this_device(device) {

```

Rather well documented:

BatchNormOptions:

https://pytorch.org/cppdocs/api/structtorch_1_1nn_1_1_batch_norm_options.html#exhale-struct-structtorch-1-1nn-1-1-batch-norm-options

BatchNorm <-- BatchNormImpl:

https://pytorch.org/cppdocs/api/classtorch_1_1nn_1_1_batch_norm_impl.html#classtorch_1_1nn_1_1_batch_norm_impl

```
    register_module("conv1", conv1);
    register_module("conv1_bn", conv1_bn);
    register_module("conv2", conv2);
    register_module("conv2_bn", conv2_bn);
    register_module("conv3", conv3);
    register_module("conv3_bn", conv3_bn);
    register_module("conv4", conv4);
    register_module("conv4_bn", conv4_bn);
    register_module("conv5", conv5);
    register_module("conv5_bn", conv5_bn);
    register_module("conv6", conv6);
    register_module("conv6_bn", conv6_bn);

    register_module("fc", fc);
    // to(this_device);
}

```

More on creating models:

<https://pytorch.org/cppdocs/frontend.html#end-to-end-example>

Forward pass

```
torch::Tensor forward(torch::Tensor x) {
```

```
    batch_size = x.size(0);
```

```
    x = x.reshape({1, 1, 1024, 1});
```

```
    x = torch::relu(conv1->forward(x));
```

```
    x = torch::max_pool2d(conv1_bn->forward(x), {2, 1});
```

```
    std::vector<torch::Tensor> cc1({x, torch::zeros({batch_size, x.size(1), 1, 1}, this_device)});  
    torch::Tensor y1 = at::cat(at::TensorList(cc1), 2);
```

```
    x = torch::relu(conv2->forward(y1));
```

```
    x = torch::max_pool2d(conv2_bn->forward(x), {2, 1});
```

```
    std::vector<torch::Tensor> cc2({x, torch::zeros({batch_size, x.size(1), 1, 1}, this_device)});  
    torch::Tensor y2 = at::cat(at::TensorList(cc2), 2);
```

```
    x = torch::relu(conv3->forward(y2));
```

```
    x = torch::max_pool2d(conv3_bn->forward(x), {2, 1});
```

```
    std::vector<torch::Tensor> cc3({x, torch::zeros({batch_size, x.size(1), 1, 1}, this_device)});  
    torch::Tensor y3 = at::cat(at::TensorList(cc3), 2);
```

```
    x = torch::relu(conv4->forward(y3));
```

```
    x = torch::max_pool2d(conv4_bn->forward(x), {2, 1});
```

```
    std::vector<torch::Tensor> cc4({x, torch::zeros({batch_size, x.size(1), 1, 1}, this_device)});  
    torch::Tensor y4 = at::cat(at::TensorList(cc4), 2);
```

```
    x = torch::relu(conv5->forward(y4));
```

```
    x = torch::max_pool2d(conv5_bn->forward(x), {2, 1});
```

```
    std::vector<torch::Tensor> cc5({x, torch::zeros({batch_size, x.size(1), 1, 1}, this_device)});  
    torch::Tensor y5 = at::cat(at::TensorList(cc5), 2);
```

```
    x = torch::relu(conv6->forward(y5));
```

```
    x = torch::max_pool2d(conv6_bn->forward(x), {2, 1});
```

```
    x = x.squeeze(3);
```

```
    x = x.permute({0, 2, 1});
```

```
    x = x.flatten(1);
```

```
    x = torch::sigmoid(fc->forward(x));
```

```
    return x;
```

```
}
```

These are documented poorly,
but are essential for complicated
models

Parsing JIT module to our model

In model definition:

```
torch::nn::Conv2d conv1{nullptr} ...
```

Modules do not contain implementations:

https://pytorch.org/cppdocs/api/classtorch_1_1nn_1_1_batch_norm.html#exhale-class-classtorch-1-1nn-1-1-batch-norm

They are derived from ModuleHolder class

Module-specific loading procedure:

```
void copy_conv2d(torch::nn::Conv2d &module, std::shared_ptr<torch::jit::script::Module> jit_module){
```

```
    module->weight.set_requires_grad(false);
```

Accessing module implementation

```
    module->weight.copy_(jit_module->get_parameter("weight"));
```

```
    module->bias.set_requires_grad(false);
```

```
    module->bias.copy_(jit_module->get_parameter("bias"));
```

Accessing jit module parameters, poorly documented:

https://pytorch.org/cppdocs/api/structtorch_1_1jit_1_1script_1_1_module.html#exhale-struct-structtorch-1-1jit-1-1script-1-1-module

```
}
```

Loading from jit script:

```
std::shared_ptr<torch::jit::script::Module> module = torch::jit::load(module_path);
```

```
copy_conv2d(net->conv1, module->get_module("conv1"));
```

JIT where do module names come from

Python code:

```
class CrepePytorchFull(nn.Module):
    def __init__(self, capacity=32):
        super(CrepePytorchFull, self).__init__()
        self.conv1 = nn.Conv2d(1, 32*capacity, ...)
        self.conv1_relu = nn.ReLU()
        self.conv1_BN = nn.BatchNorm2d(32*capacity)
        self.conv1_maxpool = nn.MaxPool2d(...)
        self.conv1_dropout = nn.Dropout(0.25)
    def forward(self, x):
        batch_size = x.size(0)
        sample_size = x.size(1)
        x = x.reshape(batch_size, 1, 1024, 1)
        x = self.conv1(x)
        x = self.conv1_relu(x)
        x = self.conv1_BN(x)
        x = self.conv1_maxpool(x)
        x = self.conv1_dropout(x)
```

JIT trace:

```
graph(%x.1 : Float(1, 1024)
      %1 : Float(512, 1, 512, 1)
      %2 : Float(512)
      %weight.1 : Float(512)
      %bias.1 : Float(512)
      %running_mean.1 : Float(512)
      %running_var.1 : Float(512))
```

Parameters names

```
...

```

```
...

```

```
...

```

Modules names

Compiling with CMake

CMake:

```
cmake_minimum_required(VERSION 3.0 FATAL_ERROR)

project(custom_ops)

find_package(Torch REQUIRED)

add_executable(example-app example-app.cpp)

target_link_libraries(example-app "${TORCH_LIBRARIES}")

set_property(TARGET example-app PROPERTY CXX_STANDARD 11)
```

```
mkdir build
```

```
cd build
```

```
cmake -DCMAKE_PREFIX_PATH=/path/to/libtorch ..
```

```
make
```

Example from https://pytorch.org/tutorials/advanced/cpp_export.html

When compiling without CUDA it uses -Wl -as-needed option for CUDA-dependent libraries.

Compiling using SConstruct

```
# SConstruct
import os
Import('env')
```

```
torch_rlibdirs = [
    "/home/lupoglaz/Projects/godot/modules/pitchdetector/libtorch/lib",
    "/usr/local/cuda-10.0/lib64"
]
```

Runtime library directories

```
torch_libdirs = [
    "/home/lupoglaz/Projects/godot/modules/pitchdetector/libtorch/lib",
    "/usr/local/cuda-10.0/lib64",
    "/home/lupoglaz/Projects/godot/modules/pitchdetector/build/Audio",
    "/usr/local/cuda-10.0/lib64",
    "/usr/lib/x86_64-linux-gnu",
    "/usr/local/cuda/lib64",
]
```

```
torch_libnames = [ "PITCH", "asound", "m", "torch", "caffe2", "caffe2_gpu", "c10", "c10_cuda",
    "cufft", "curand", "cudnn", "culibos", "cublas", "cuda", "nvrtc",
    "nvToolsExt", "cudart_static", "pthread", "dl", "rt", "X11"]
```

Libraries that we need

```
env_sum = env.Clone()
env_sum.Append(LINKFLAGS=['-Wl,--no-undefined', '-Wl,--no-as-needed'])
env_sum.Append(CPPPATH=torch_include)
env_sum.Append(LIBPATH=torch_libdirs)
env_sum.Append(RPATH=torch_rlibdirs)
env_sum.Append(LIBS=torch_libnames)
env_sum.Append(CPPDEFINES={'-D_GLIBCXX_USE_CXX11_ABI' : 0 })
env_sum.Append(CXXFLAGS=['-O2', '-Wall', '-std=c++11'])
```

Important option

```
src_list = ["pitchdetector.cpp", "register_types.cpp", "Audio/cAlsaIO.cpp", "Audio/cNNPitchDetector.cpp"]
env_sum.add_source_files(env.modules_sources, src_list)
```

```
env.Append(LINKFLAGS=['-Wl,--no-undefined', '-Wl,--no-as-needed'])
env.Append(LIBPATH=torch_libdirs)
env.Append(RPATH=torch_rlibdirs)
env.Append(LIBS=torch_libnames)
```

Linking libraries to the global environment

Demo

Thanks